

EIDOLON

Can you let an AI agent act on your behalf?

Only if it's bounded, restrained,
revocable, and fully attributable.

Provable delegated agency.

A digital twin that decides the way you would — and can prove it.

THE PROBLEM

Today's agents pick one bad option.

Over-permissioned

Give it your tools and hope. One bad call — a wrong email, a deleted database.

= unsafe

Locked down

Approve every step. The human is the bottleneck again — no leverage.

= useless

The missing piece isn't a smarter model.

It's a **governance layer** for authority.

Delegate a bounded slice of your authority.

- Fidelity & authority are separate.
“Would you act?” and “is the twin permitted?” both must pass.
- Authority only attenuates, never widens.
Signed, chained credentials — a twin can't exceed you.
- Restraint is the product.
It escalates when unsure or out of scope — and hands back.
- Every action is attributable.
Attested on a consensus ledger — the whole session replays.

Revocable. Cryptographic. Auditable.

Pull the authority back mid-session; the next action is denied.

THE GATE

One choke point. Five honest outcomes.

Every candidate action resolves through KAIROS, in a locked order.

DENY

not authorized — invalid or revoked credential

ESCALATE

out of scope, low confidence — hand back to you

DRAFT

prepared in your voice — held for approval

NOTIFY

recoverable action — done, you're informed

ACT


fully authorized — executed autonomously

Attest-then-act: no side effect without a prior attestation.

THE INTEGRATION

The authority layer for any MCP agent.

SAGE became the memory layer. EIDOLON is the authority layer.



Point your agent at the gateway.

Zero code changes. Every tool call is now governed & attested.

✓ read status → runs

⚠ delete prod DB → escalated

🛑 draft email → held

✗ out-of-scope host → denied

Works with Hermes · Claude Code · OpenClaw · Raptor · Cursor.

WHAT'S REAL

Not a slide deck. A working system.

- ✓ Two invariants, property-tested
default-deny · no unattested action
- ✓ Injection-resistant by design
authority is memory-blind — prompts can't widen it
- ✓ Verified on a live consensus ledger
attest → replay, byte-identical
- ✓ Governs real agents over MCP
continuity twin · governed red-teamer · ops tools

Building agents that act for people?

Let's talk about the authority layer. Follow + DM · repo in comments.